

ARMv8 - 64-bit Architecture

First look at the arm64
architecture and assembly
language


Assembly Language and Architecture

- Assembly-language programs consist of instructions that manipulate processor hardware elements
 - Goal of manipulations:
implement high-level-language program behaviors
- Processor designs determine what instructions are available
- Understanding assembly language requires understanding the processor architecture

Computer Organization – Programmer's View

- Fundamental components:

- CPU
- Memory
- I/O ("everything else")



- Possible connection schemes:

- Single Controller (right)
- Direct Memory Access (DMA) (below)



Memory Organization

- The memory subsystem is a 1-dimensional array of storage locations

- Each location has a unique address (or "array index")
- Each location is made up of m bits ($m \geq 1$)
- Location's contents is any one of 2^m bit patterns

- Memory can hold:

- Programs: bit patterns represent instructions
- Data: bit patterns represent numbers, characters, etc.

- Volatile contents: R/W memory, or RAM


- DRAM is common for main memory

- Nonvolatile (unchangeable): Read-Only, or ROM

- PROM, EEPROM, Flash – erasable/reprogrammable ROM


How a program occupies memory

- When a program runs it occupies a block of RAM



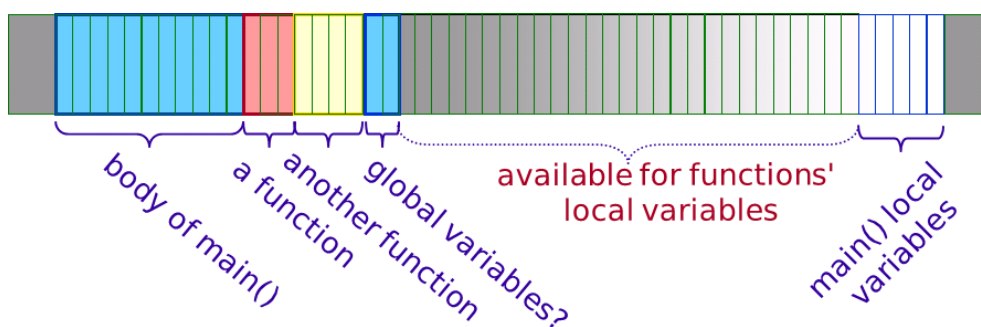
Your program's memory block

- The block is divided into instructions, data, and stack segments, each in part of the memory




Dividing up the program memory

- Each function, including `main()`, has its own unchanging portion of the instruction memory
 - Any global variables occupy memory similar to instruction memory.
- Function-local variables exist in the stack only while the function is executing



the Central Processing Unit

- a.k.a. CPU or Processor
 - (or microprocessor)
- Reads program instructions
- Accesses program data
- Performs numerical and logical operations
- von Neumann model
 - basis for most modern computers



John von Neumann's IAS Processor



- Developed at the Institute for Advanced Studies, Princeton, NJ
- Single, shared memory for program instructions and data

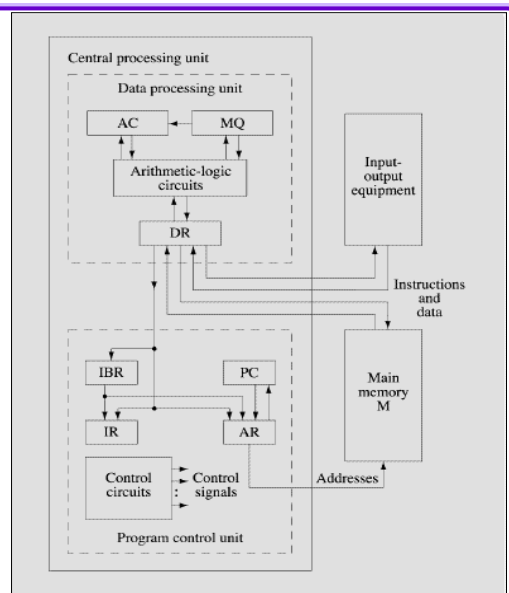
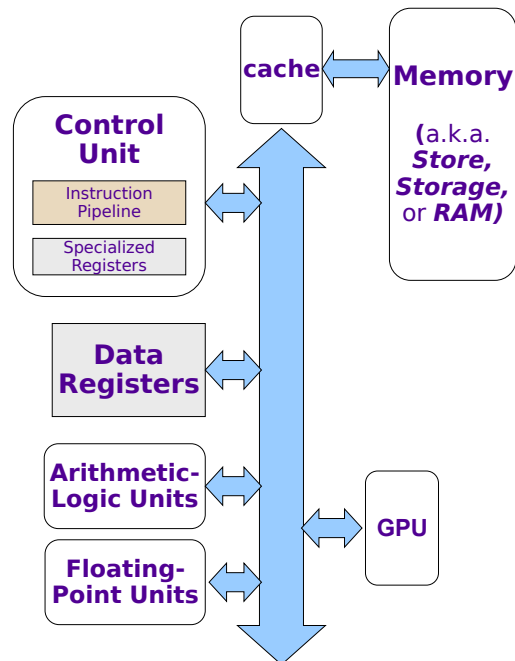


Figure 1

Structure of von Neumann's IAS processor. Reprinted with the permission of the McGraw-Hill Book Company from page 23 of the second edition of *Computer Architecture and Organization* by John P. Hayes, published by the McGraw-Hill Book Company, copyright 1988.

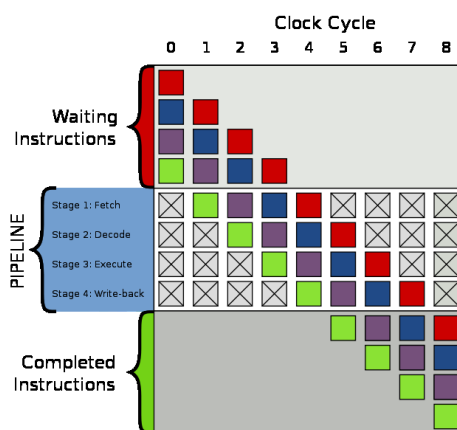

Modern CPU elements

- Control Unit
 - Instruction pipeline
 - Specialized registers
 - » contain execution state
 - » control program flow
- General-Purpose Registers
 - Hold data values for ALU, FPU
- ALU, FPU, GPU
 - Do the computations
- Cache
 - Improves memory accesses
 - Generally not visible to programmer
- Internal buses
 - Connect things



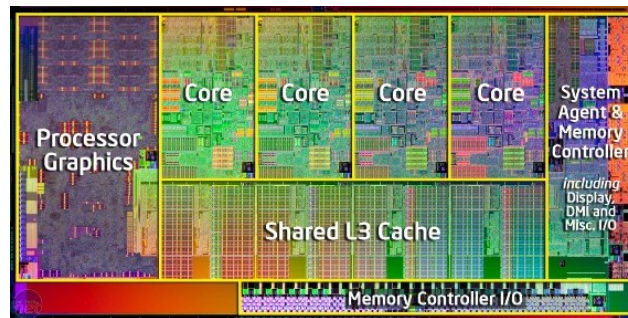

Pipelined Designs

- Separate circuits perform each stage of the fetch-execute cycle
- In a *pipelined* design the separate circuits work simultaneously, on sequential instructions

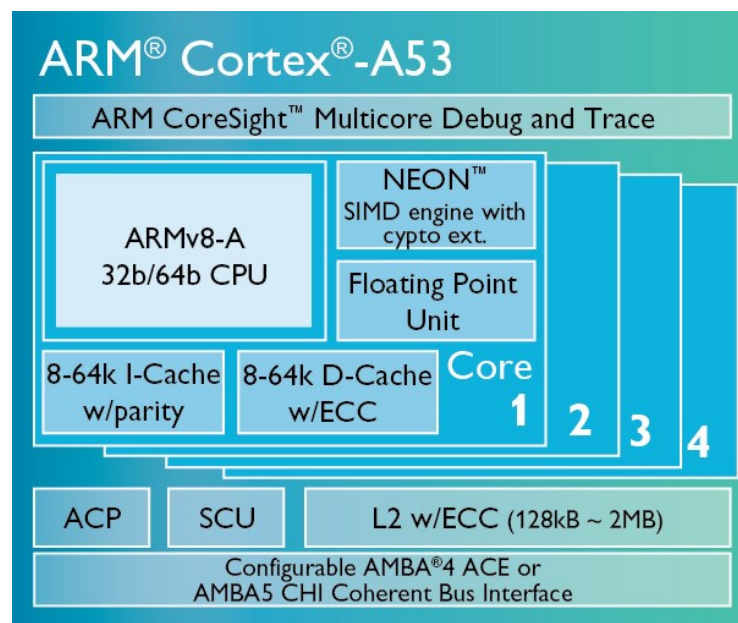


Multi-core Processors

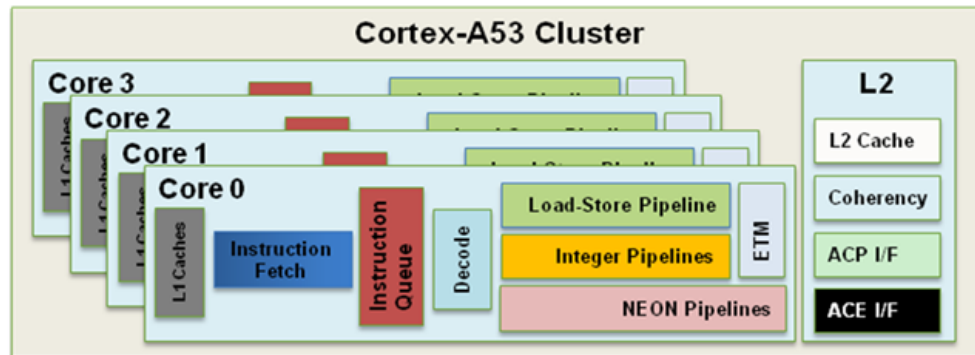
- Each CPU is a "core"
- One chip holds 2, 4, 8, ... cores
- Cores share memory
- Cores may have local L1 cache, may share L2 and L3 cache
- A CPU with a GPU is a form of *heterogeneous* multicore processor



Raspberry Pi's Processor Block Diagram



RPi Instruction Pipeline



The ARM Architectural Models

- Aarch32 (armhf, armel)
 - Original models, optional floating-point
 - Provides 32-bit registers and instructions
- T32 (Thumb2)
 - 16-bit instructions expand to 32 bits on use
 - Compatible with Aarch32
- Aarch64 (arm64)
 - 64-bit registers, some new instructions
 - Floating-point, SIMD hardware standard
 - Context switch between Aarch64 , Aarch32

Aarch64 General-Purpose Register Set

- r0 - r7
 - Arguments, return values
- r8
 - Syscall number / Returned-struct pointer
- r9 - r15
 - Temporary
- r16 - r18
 - Intra-procedure/platform
 - Avoid using these
- r19 - r28
 - Callee preserves these
- r29, r30
 - Frame, link registers:
 - Don't use for general purpose

Register names

- X-names
 - 64-bit versions
- W-names
 - Lower 32 bits

Syscall number
/ Return pointer

Frame pointer

Link register

Zero register
/ stack pointer

X0	W0	r0
X1	W1	r1
X2	W2	r2
X3	W3	r3
X4	W4	r4
X5	W5	r5
X6	W6	r6
X7	W7	r7
X8	W8	r8
X9	W9	r9
X10	W10	r10
X11	W11	r11
X12	W12	r12
X13	W13	r13
X14	W14	r14
X15	W15	r15
X16	W16	r16
X17	W17	r17
X18	W18	r18
X19	W19	r19
X20	W20	r20
X21	W21	r21
X22	W22	r22
X23	W23	r23
X24	W24	r24
X25	W25	r25
X26	W26	r26
X27	W27	r27
X28	W28	r28
X29	W29	r29
X30	W30	r30
Xzr	Wzr	r2r

← 32 bits →
← 64 bits →

Aarch64 Special Registers


Zero Register
Program Counter
Stack Pointer
Saved Program Status register
Exception Link register

XZR	WZR
PC	
SP	WSP
SPSR_EL1 / SPSR_EL2 / SPSR_EL3	
ELR_EL1 / ELR_EL2 / ELR_EL3	

- Zero Register
 - reads always return 0
 - writes are always discarded
- PC
 - no explicit access
- SP, SPSR, EL registers:
 - Separate set for each Exception level EL0 - EL 3
- SPSR
 - includes NZCV condition-code flags
- EL
 - holds exception return address
- Some instructions treat ZR or SP as "r31"

Aarch64 VFP/SIMD (NEON) Register Set

- V0 - V7
 - Arguments, return values
- V8 - V15
 - Callee-saved
(lower 64 bits only)
- V16 - V31
 - Spare temporary registers
 - » Not in Aarch32
- D0 - D31
 - Double-precision registers
 - » D0 - D15 also in Aarch32
- S0 - S31
 - Single-precision registers
 - » S0 - S15 also in Aarch32



Exception Levels

- ARMv8 security model
- "Userland" operates at Level 0 (L0)
 - Least privilege
 - Library Functions
- Operating system operates at L1
 - System calls
- Virtualization runs at L2
- Hardware has highest privilege level, L3

