

Linux Kernel Overview

Manolis Marazakis


Michail Flouris

FORTH-ICS (CARV Lab)

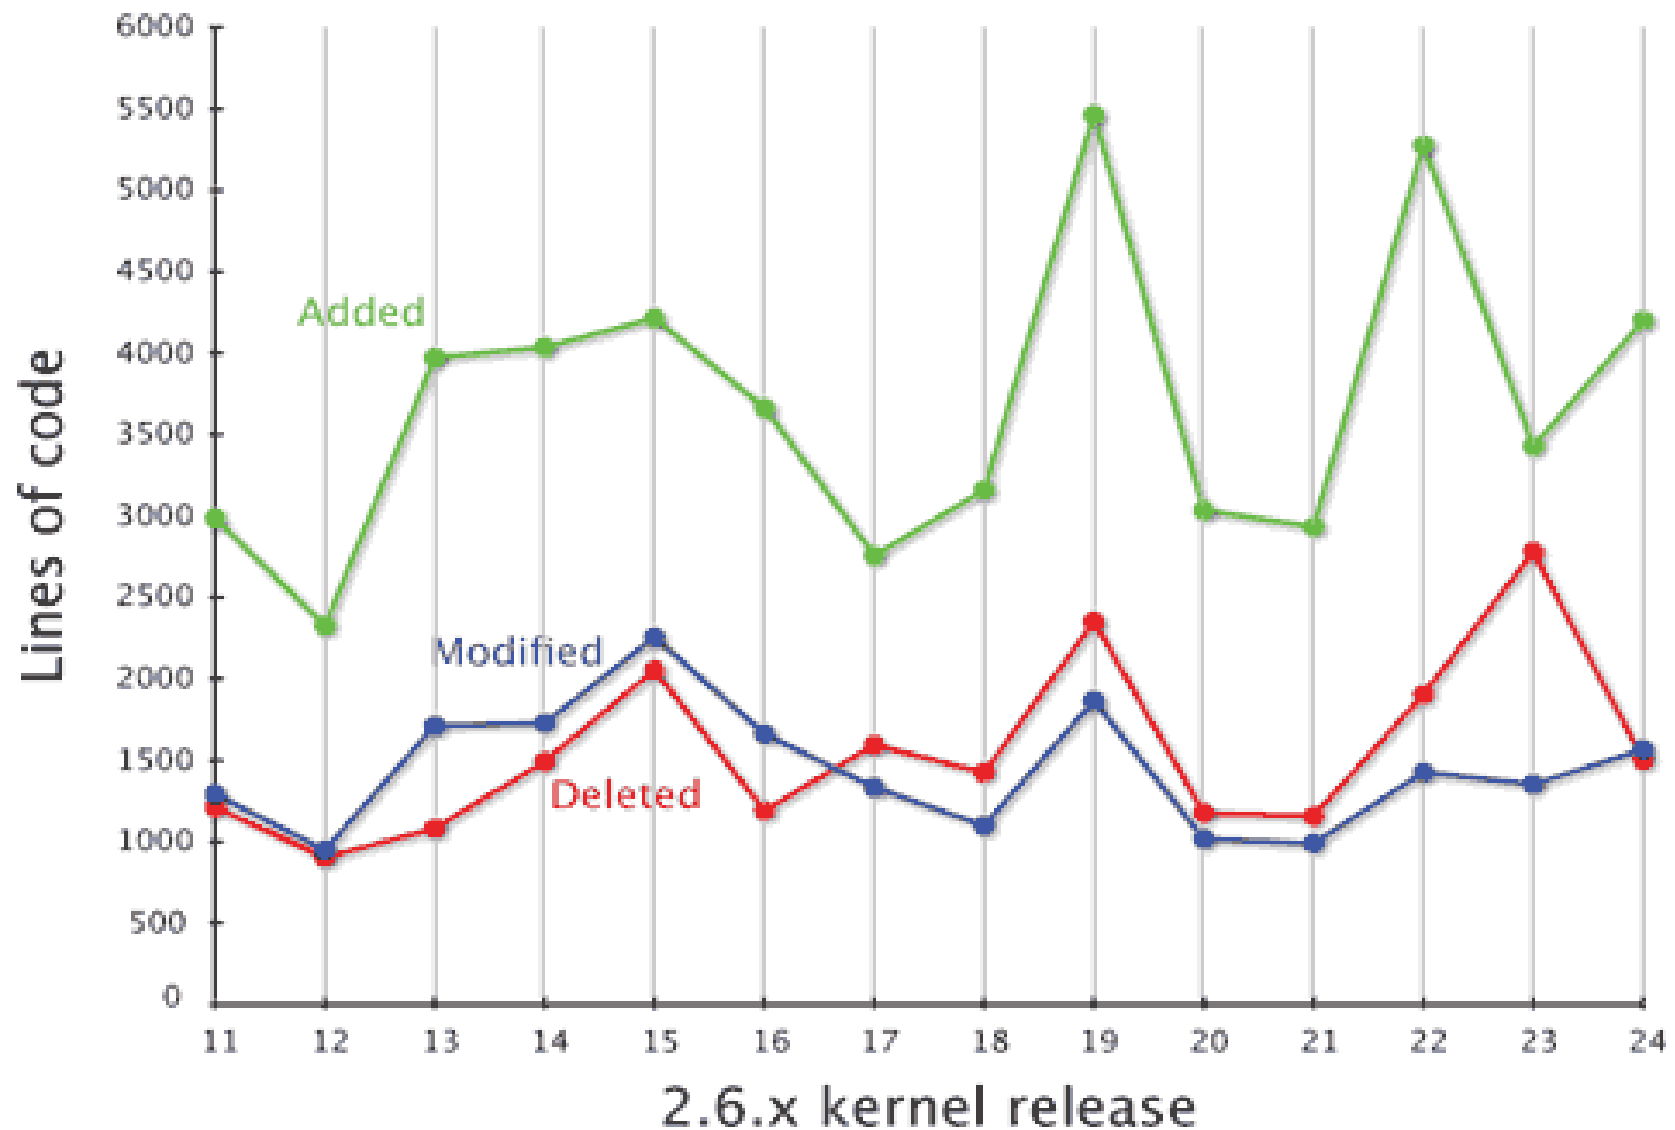
References

- J. Corbet, A. Rubini, G. Kroah-Hartman: Linux Device Drivers (O'Reilly, 3rd edition)
- R. Love: Linux Kernel Development (Novell, 2nd edition)
- C. Hallinan: Embedded Linux Primer: A Practical Real-World Approach (Prentice Hall)
- S. Venkateswaran: Essential Linux Device Drivers (Prentice Hall)
- LWN.net
- LXR – Linux Cross Reference: <http://lxr.linux.no/>
- <http://www.linuxfoundation.org/>
- LKML.org

Linux Kernel: size



Linux Kernel: rate of change



Linux Internals

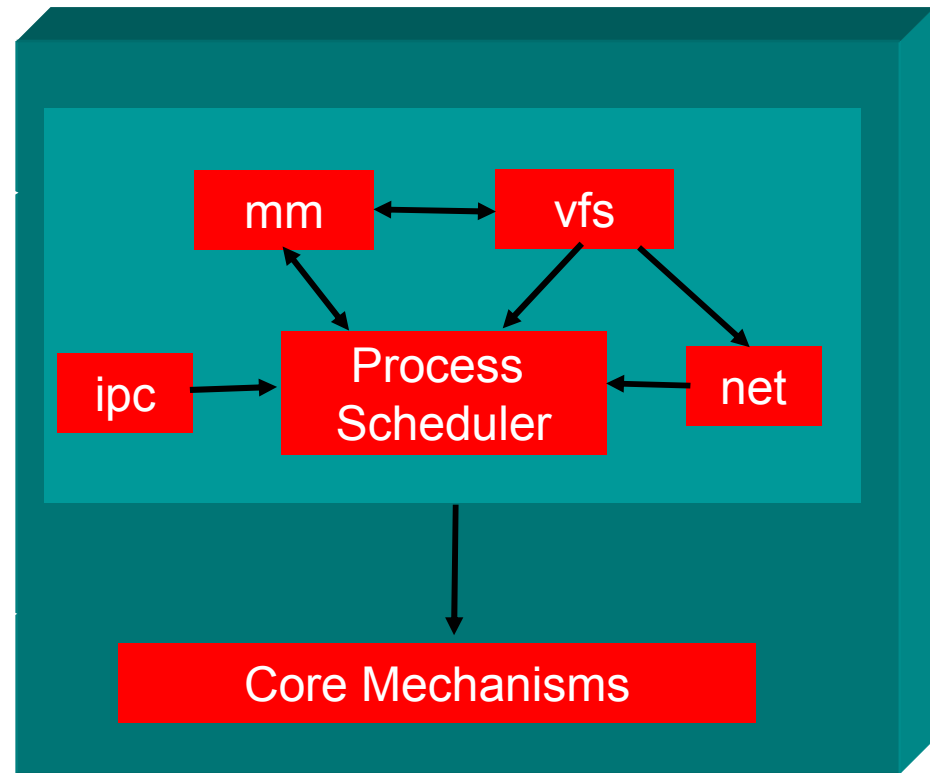
Total LOC for 2.6.18 kernel: 5,499,231

Modules

drivers: 2,748,214 LOC

Device
Drivers

arch: 920,662 (24 variations:
powerpc, arm, x86_64, i386, ia64,
sparc, sparc64, mips, sh, sh64,
m68k, s390, ...)



core kernel: 47,115

init: 2,215

ipc: 4,681 | mm: 24,618 | net: 336,302

Device Drivers

- Black boxes to hide details of hardware devices
- Use standardized calls
 - Independent of the specific driver
- Main role
 - Map standard calls to device-specific operations
- Can be developed separately from the rest of the kernel
 - Plugged in at runtime when needed

Role of a Device Driver

- Implements *mechanisms* to access the hardware
 - E.g., show a disk as an array of data blocks
- Does not enforce particular *policies* on users
 - Examples:
 - Who may access the drive?
 - Is the drive accessed via a file system?
 - May users mount file systems on the drive?

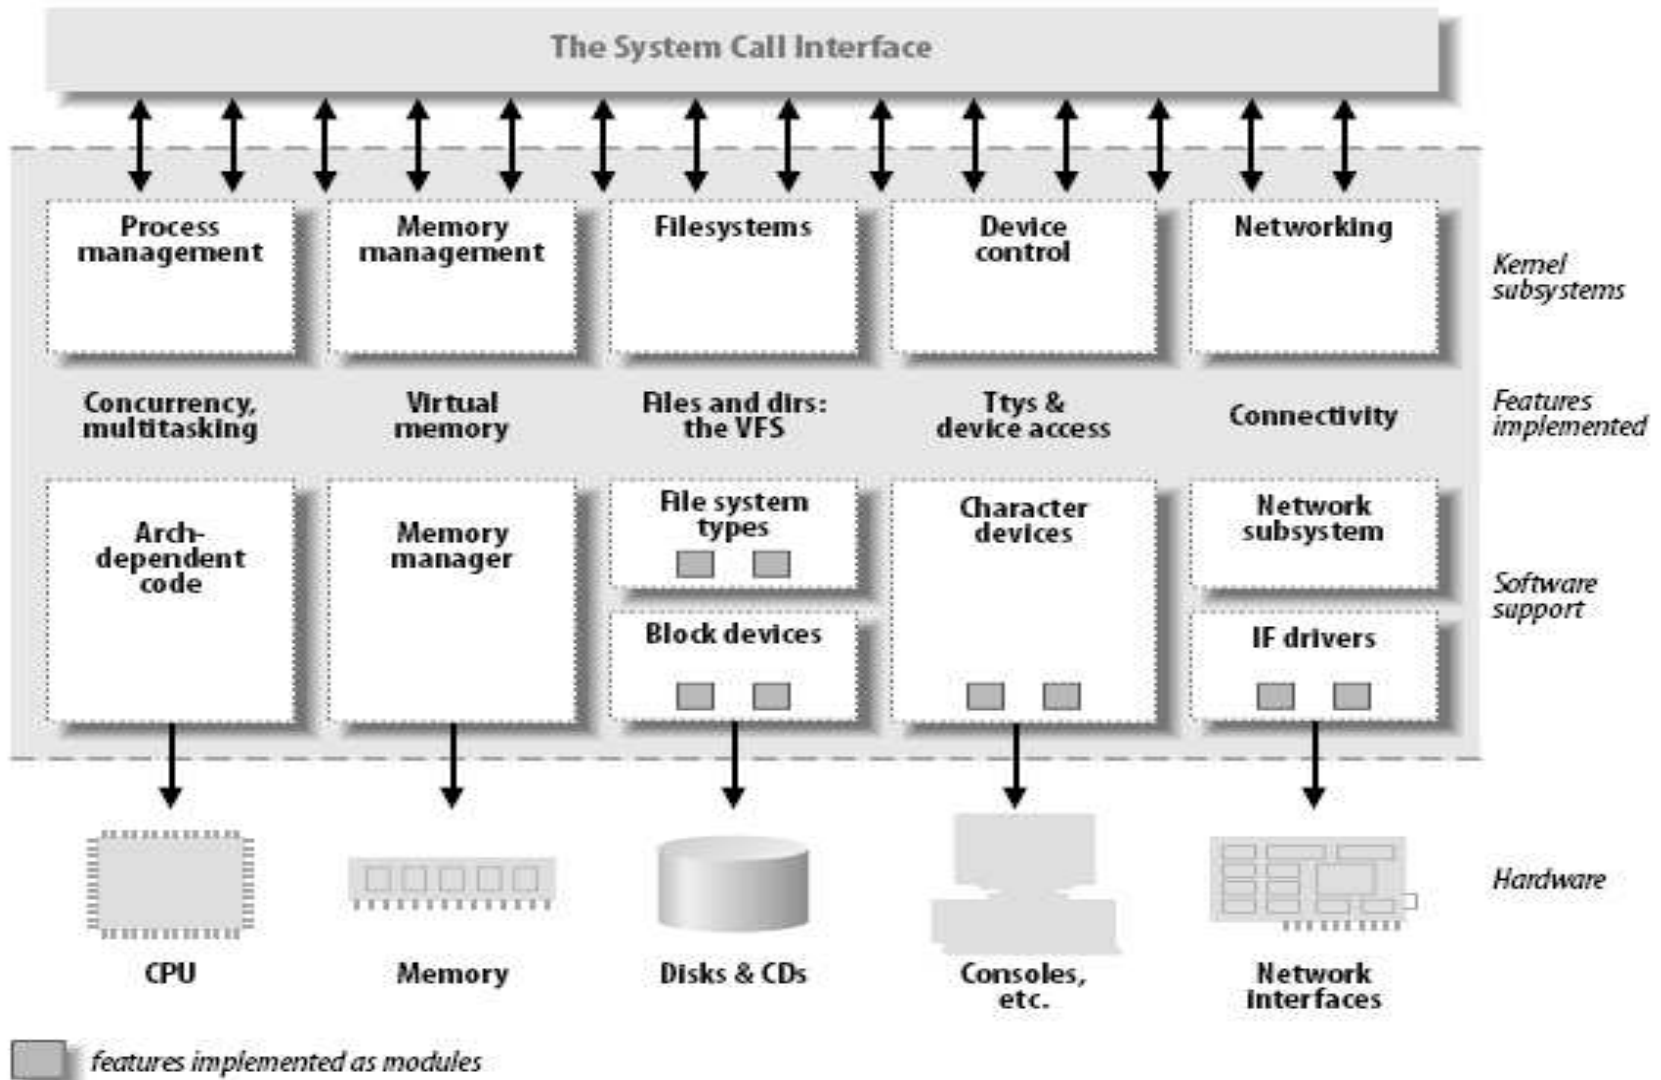
Kernel Subsystems

- Process management
 - Creates, destroys processes
 - Supports communication among processes
 - Signals, pipes, etc.
 - Schedules how processes share the CPU
- Memory management
 - Virtual addressing

Kernel Subsystems

- File systems
 - Everything in UNIX can be treated as a file
 - Linux supports multiple file systems
- Device control
 - Every system operation maps to a physical device
 - Few exceptions: CPU, memory, etc.
- Networking
 - Handles packets
 - Handles routing and network address resolution issues

Kernel Subsystems



Loadable Kernel Modules

- The ability to add and remove kernel features at runtime
- Each unit of extension is called a *module*
- Use *insmod* program to add a kernel module
- Use *rmmod* program to remove a kernel module

Role of a Module

- Extend kernel functionality
 - Modularized code running in kernel space
- Example: filesystems (xfs, jffs2, fuse, ...)

Classes of Devices and Modules

- Character devices
- Block devices
- Network devices
- Others: USB, SCSI, FireWire, I2C, MTD

Character Devices

- Abstraction: a stream of bytes
 - Examples
 - Text console (`/dev/console`)
 - Serial ports (`/dev/ttyS0`)
 - Usually supports **open**, **close**, **read**, **write**
 - Accessed sequentially (in most cases)
 - Might not support file seeks
 - Exception: frame grabbers
 - Can access acquired image using **mmap** or **lseek**

Block Devices

- Abstraction: array of storage blocks
- However, applications can access a block device in bytes
 - Block and char devices differ only at the kernel level
 - A block device can host a file system

Network Devices

- Abstraction: data packets
- Send and receive packets
 - Do not know about individual connections
- Have unique names (e.g., **eth0**)
 - Not in the file system
 - Support protocols and streams related to packet transmission (i.e., no **read** and **write**)

Filesystem Modules

- Software drivers, not device drivers
- Serve as a layer between user API and block devices
- Intended to be device-independent

HelloWorld module

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("GPL");
```

```
static __init int hello_init(void) {
    printk(KERN_ALERT "Hello, world\n");
    return 0;
}
```

```
static __exit void hello_exit(void) {
    printk(KERN_ALERT "Goodbye, cruel world\n");
}
```

```
module_init(hello_init);
module_exit(hello_exit);
```

% cat Makefile

obj-m := module.o

module-objs := file1.o file2.o

% make -C /usr/src/linux M=`pwd` modules

Kernel Modules vs. Applications

- Applications
 - Can access various functions in user-level libraries (e.g., `printf` in C library)
- Kernel modules
 - No user-level libraries
 - `printk` is defined within the kernel
 - Exported to modules
 - Should include only header files defined within the kernel source tree

Threads/Processes

- Thread: A sequential execution stream
- Address space: Chunks of memory and everything needed to run a program
- Process: An address space + thread(s)

User Space vs Kernel Space

- Kernel modules run in *kernel space*
 - Execute in the *supervisor mode*
 - Everything is allowed
 - Share the same address space
- Applications run in *user space*
 - Execute in the *user mode*
 - Restricted access to hardware
 - Each has its own address space

System Calls

- *System calls* allow processes running at the *user mode* to access kernel functions that run under the *kernel mode*
- Prevent processes from doing bad things, such as
 - Halting the entire operating system
 - Modifying the MBR

User Level Drivers

- + Fast development
- + C library support
- + Conventional debugger
- + Fault isolation
- + Portability
- Interrupts not available
- Privileged access required for direct memory access
- Poor performance

Hardware Interrupts

- Can suspend user-level processes
 - Transfers execution from user space to kernel space
 - Interrupts are handled by separate threads
 - Not related to any user-level processes
 - Asynchronous

Concurrency in the Kernel

- Sources of concurrency
 - Hardware interrupts
 - Kernel timers
 - Multiple CPUs
 - Preemption

Handling Concurrency

- Kernel code needs to be *reentrant*
 - Capable of running in more than one thread execution context at the time
 - Prevent corruption of shared data
 - Avoid *race conditions*
 - Results depend on the timing of their executions

DRAM-peek module (I)

```
#include <linux/module.h>
#include <linux/highmem.h>
#include <asm/uaccess.h>
```

```
char modname[] = "dram";
int my_major = 253;
unsigned long dram_size;
```

```
static loff_t my_llseek( struct file *file, loff_t offset, int whence );
static ssize_t my_read( struct file *file, char *buf, size_t count, loff_t *pos );
```

```
struct file_operations my_fops = {
    owner:      THIS_MODULE,
    llseek:     my_llseek,
    read:       my_read,
};
```

DRAM-peek module (II)

```
static __init int init_module(void) {
    printk( "<1>\nInstalling '%s' module ", modname );
    printk( "(major=%d)\n", my_major );
    dram_size = num_physpages * PAGE_SIZE;
    printk( "<1> ramtop=%08lX (%lu MB)\n",
            dram_size, dram_size >> 20 );
    return register_chrdev( my_major, modname, &my_fops );
}

static __exit void cleanup_module(void) {
    unregister_chrdev( my_major, modname );
    printk( "<1>Removing '%s' module\n", modname );
}
```

DRAM-peek module (III)

```
static ssize_t my_read( struct file *file, char *buf, size_t count, loff_t *pos ) {
    struct page      *pp;
    void              *from;
    int               page_number, page_indent, more;

    if ( *pos >= dram_size ) return 0;
    page_number = *pos / PAGE_SIZE;
    page_indent = *pos % PAGE_SIZE;

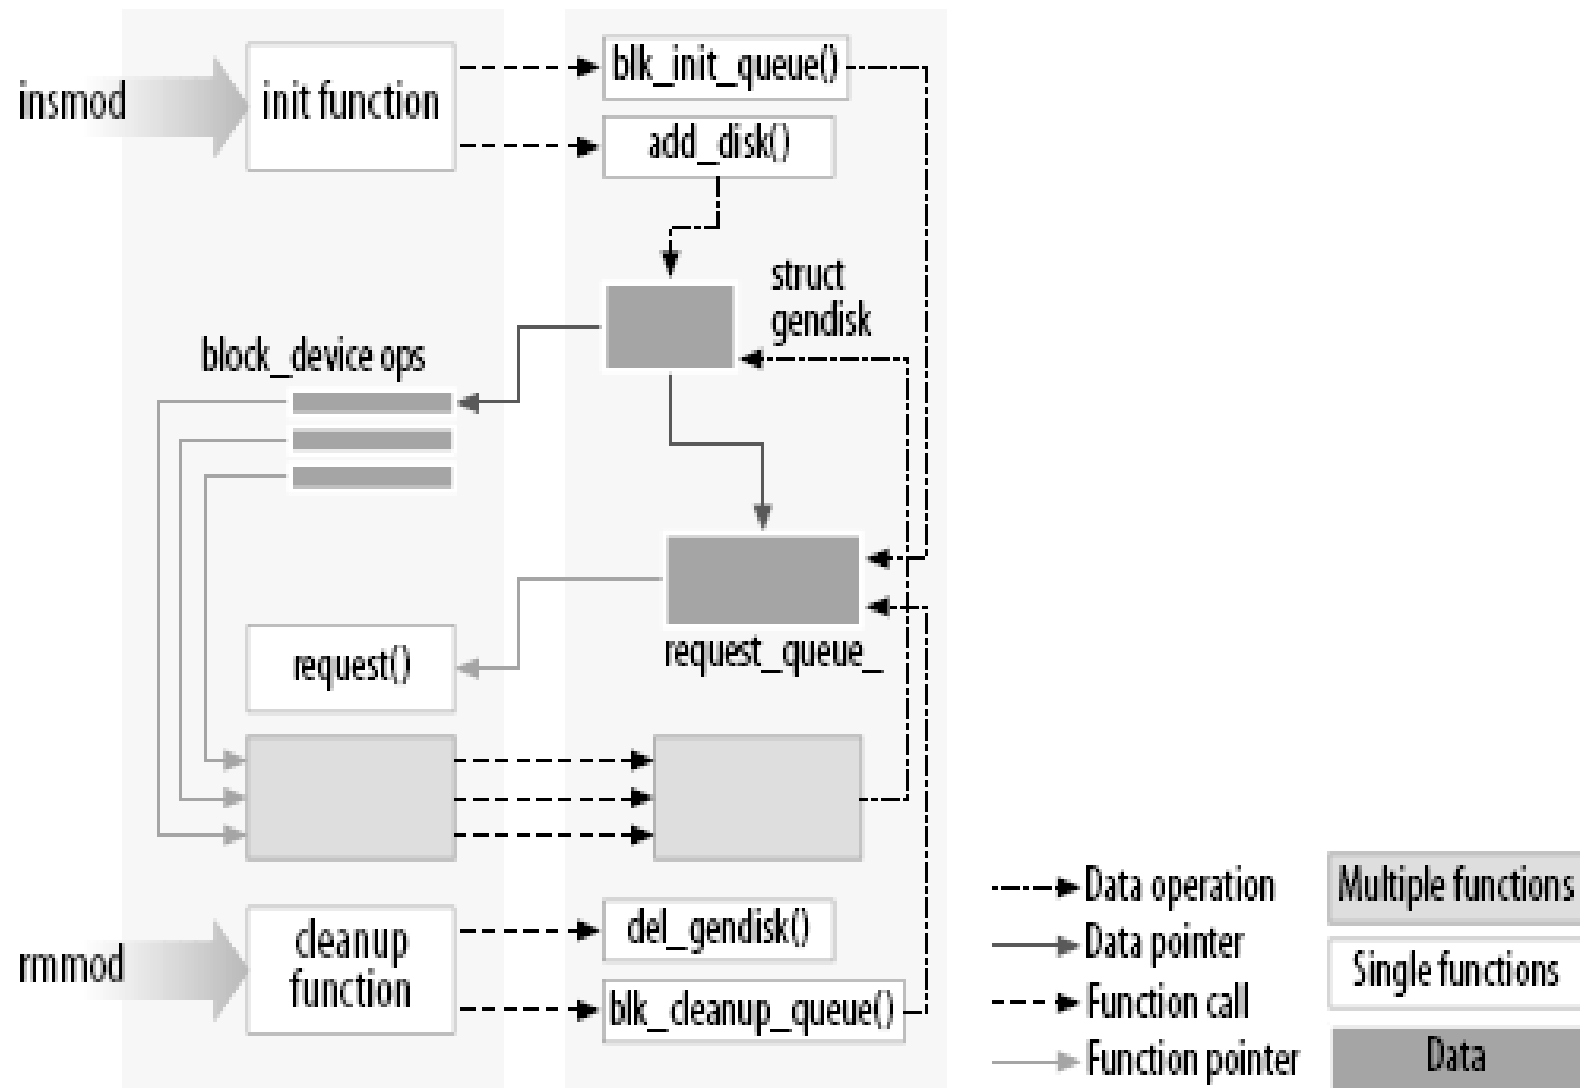
    pp = &mem_map[ page_number ];
    from = kmap( pp ) + page_indent;
    if ( page_indent + count > PAGE_SIZE ) {
        count = PAGE_SIZE - page_indent;
    }
    more = copy_to_user( buf, from, count );
    kunmap( pp );
    if ( more ) { return -EFAULT; }
    *pos += count;
    return count;
}
```

DRAM-peek module (IV)

```
static loff_t my_llseek( struct file *file, loff_t offset, int whence ) {
    loff_t    newpos = -1;
    switch( whence ) {
        case 0: newpos = offset; break;           // SEEK_SET
        case 1: newpos = file->f_pos + offset; break; // SEEK_CUR
        case 2: newpos = dram_size + offset; break; // SEEK_END
    }
    if (( newpos < 0 )||( newpos > dram_size )) return -EINVAL;
    file->f_pos = newpos;
    return    newpos;
}
```

```
MODULE("GPL");
```

Linking a Module to the Kernel



Loading and Unloading Modules

- **insmod**

- Links unresolved symbol in the module to the symbol table of the kernel
- **cat /proc/modules** to see list of currently loaded modules

- **rmmmod**

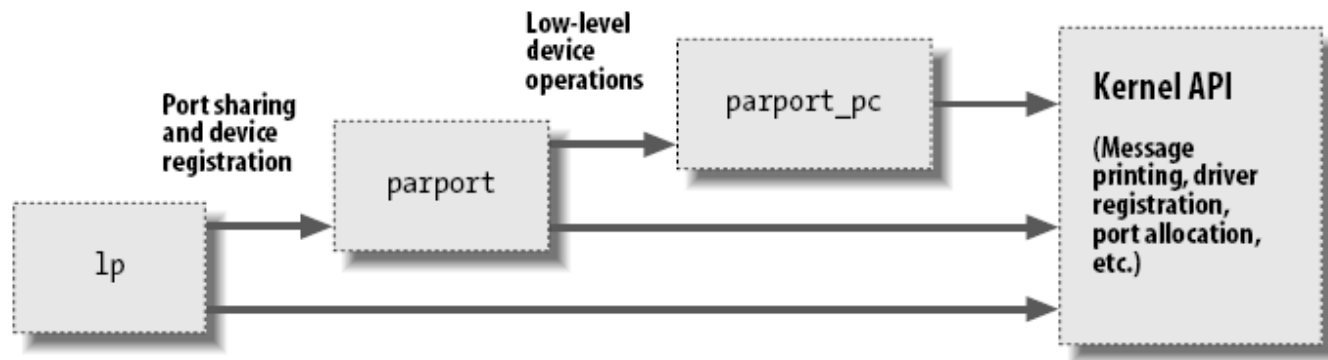
- Removes a kernel module
- Fails when the kernel believes that it is still in use
 - Or, something has gone wrong
 - Might need to reboot to remove the module

Kernel Symbol Table

- Addresses of global functions and variables
- A module can export its symbols for other modules to use
- *Module stacking*
 - E.g., MSDOS file system relies on symbols exported by the FAT module

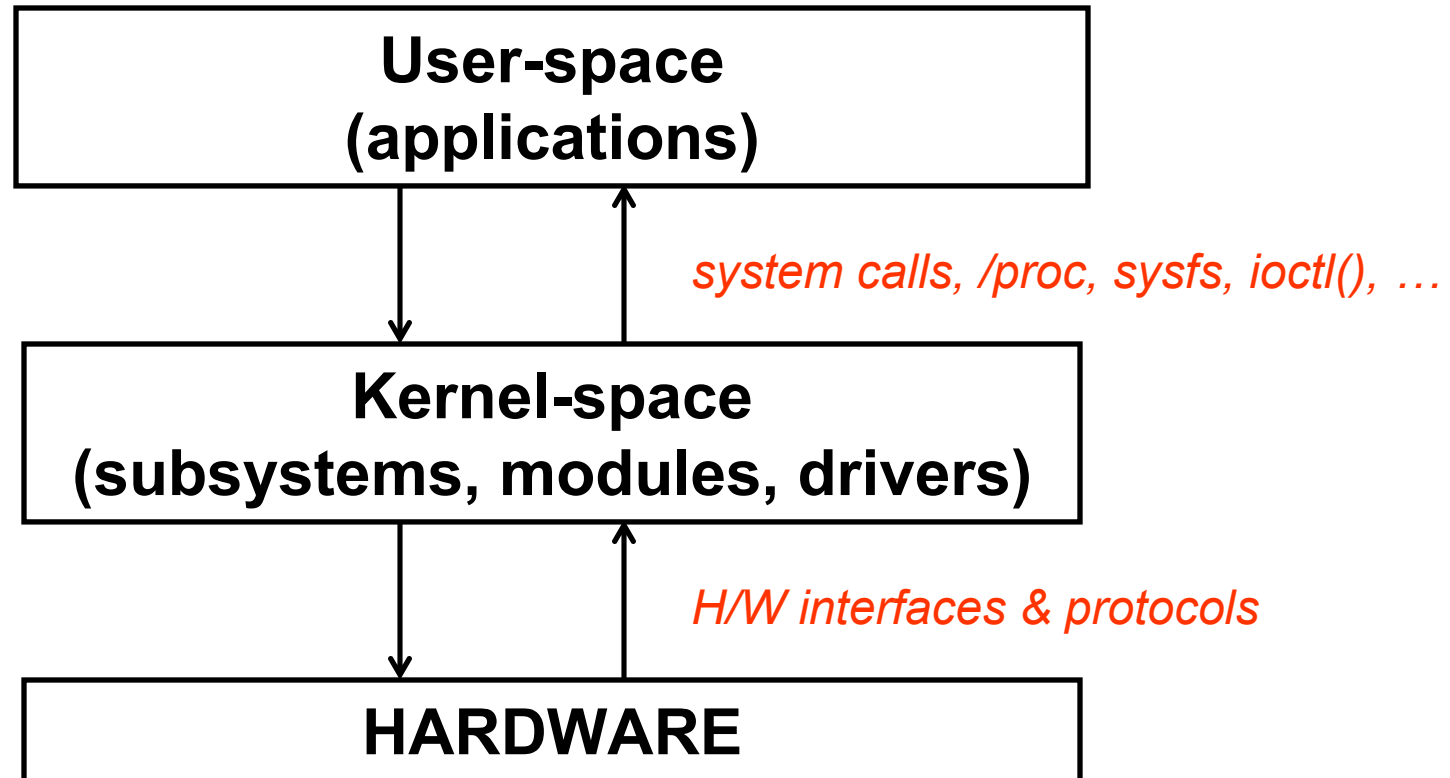
Module Stacking Example

- Stacking of parallel port driver modules



- Can use **modprobe** to load all modules required by a particular module

User-space vs Kernel-space



Device driver **events**

Device driver **interfaces**

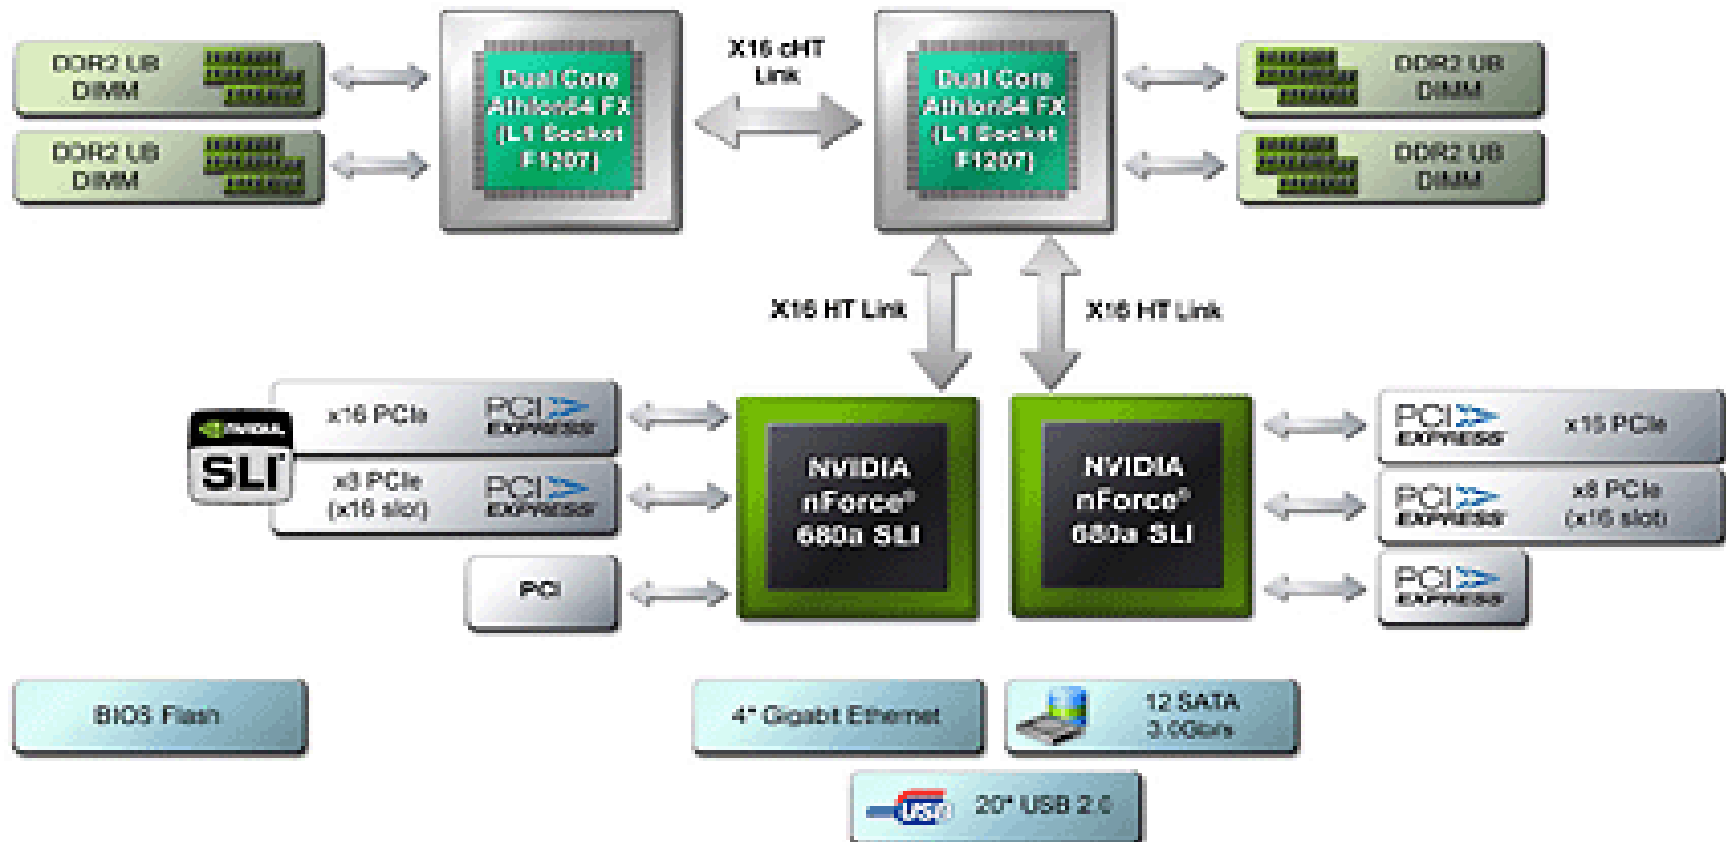
User - Kernel Space Interfaces

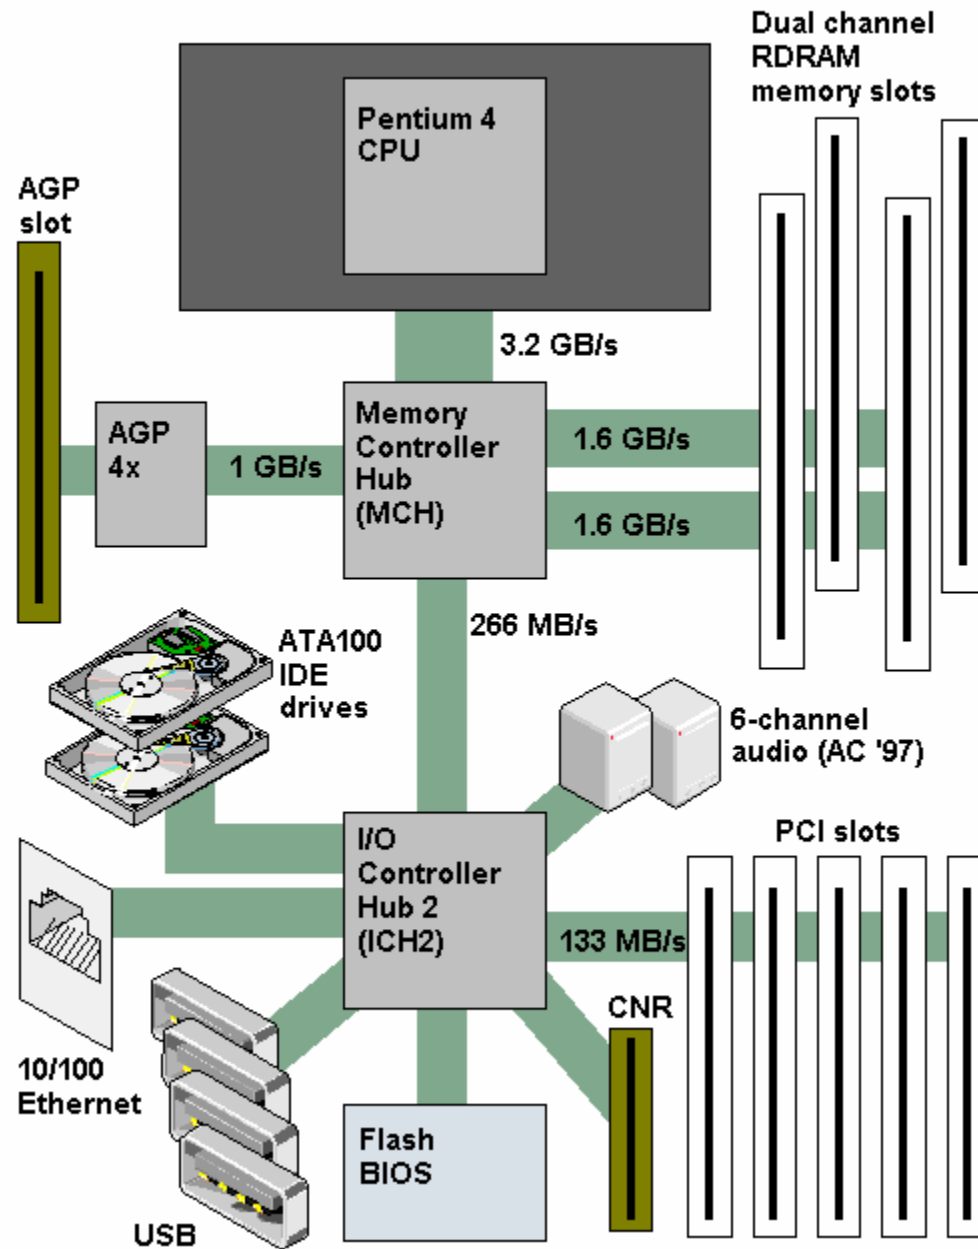
Events	User functions	Kernel functions
Load module	insmod	module_init()
Open device	fopen	file_operations: open
Close device	fread	file_operations: read
Write device	fwrite	file_operations: write
Close device	fclose	file_operations: release
Remove module	rmmod	module_exit()

Chipset (AMD/NVIDIA)



NVIDIA nForce® 680a SLI™ System Architecture
Motherboard Works with 1 or 2 CPUs





Chipset (Intel/ICH)

Chipset (Intel/MCH)

